

Introduction to Data Formats and S3

```
In [1]: import pandas as pd
import numpy as np

import boto3
import sagemaker.amazon.common as smac
```

```
In [2]: np.random.seed(5)
```

```
In [3]: # NOTE: Specify your bucket
s3_bucket_name = 'dwb-ml-sagemaker'
```

Sample DataSet

Three features x1,x2,x3 and a target variable y

```
In [4]: n = 10

x1 = np.random.random_sample(n)      # n floating point numbers between 0 and 1
x2 = np.random.randint(100,200,n)    # n integers
x3 = np.random.random_sample(n) * 10 # n floating point numbers between 0 and 10
y = np.random.randint(0,2,n)         # Response variable 0 or 1
```

```
In [5]: y
```

```
Out[5]: array([0, 0, 1, 1, 1, 1, 0, 0, 0, 1])
```

```
In [6]: df = pd.DataFrame({'x1':x1,
                           'x2':x2,
                           'x3':x3,
                           'y':y})
```

```
In [7]: df
```

```
Out[7]:
```

	x1	x2	x3	y
0	0.221993	153	2.041547	0
1	0.870732	180	1.190954	0
2	0.206719	127	8.779031	1
3	0.918611	144	5.236753	1
4	0.488411	177	4.921360	1
5	0.611744	175	7.318711	1
6	0.765908	165	0.145808	0
7	0.518418	147	0.933630	0
8	0.296801	130	8.265542	0
9	0.187721	184	8.334927	1

```
In [8]: # Write to SageMaker Notebook Instance
df.to_csv('demo_file.csv', index=False)
```

```
In [9]: # Write and Reading from S3 is just as easy
# files are referred as objects in S3.
# file name is referred as key name in S3
# Files stored in S3 are automatically replicated across 3 different availability z
# in the region where the bucket was created.

# http://boto3.readthedocs.io/en/latest/guide/s3.html
def write_to_s3(filename, bucket, key):
    with open(filename, 'rb') as f: # Read in binary mode
        return boto3.Session().resource('s3').Bucket(bucket).Object(key).upload_file(f)
```

```
In [10]: # http://boto3.readthedocs.io/en/latest/guide/s3.html
def download_from_s3(filename, bucket, key):
    with open(filename, 'wb') as f:
        return boto3.Session().resource('s3').Bucket(bucket).Object(key).download_file(f)
```

```
In [11]: write_to_s3('demo_file.csv', s3_bucket_name, 'data_format/demo_file.csv')
```

```
In [12]: download_from_s3('demo_file_from_s3.csv', s3_bucket_name, 'data_format/demo_file.csv')
```

RecordIO Format

We will use SageMaker SDK `write_numpy_to_dense_tensor()` method to create RecordIO files

Data Types: Int32, Float32, Float64

Reference: <https://github.com/aws/sagemaker-python-sdk/blob/master/src/sagemaker/amazon/common.py>

```
In [13]: df.head()
```

```
Out[13]:
```

	x1	x2	x3	y
0	0.221993	153	2.041547	0
1	0.870732	180	1.190954	0
2	0.206719	127	8.779031	1
3	0.918611	144	5.236753	1
4	0.488411	177	4.921360	1

```
In [14]: # X must be an array
X = df[['x1', 'x2', 'x3']].to_numpy()
```

```
In [15]: X
```

```
Out[15]: array([[2.21993171e-01, 1.53000000e+02, 2.04154748e+00],
 [8.70732306e-01, 1.80000000e+02, 1.19095357e+00],
 [2.06719155e-01, 1.27000000e+02, 8.77903071e+00],
 [9.18610908e-01, 1.44000000e+02, 5.23675290e+00],
 [4.88411189e-01, 1.77000000e+02, 4.92135999e+00],
 [6.11743863e-01, 1.75000000e+02, 7.31871100e+00],
 [7.65907856e-01, 1.65000000e+02, 1.45807511e-01],
 [5.18417988e-01, 1.47000000e+02, 9.33630336e-01],
 [2.96800502e-01, 1.30000000e+02, 8.26554249e+00],
 [1.87721229e-01, 1.84000000e+02, 8.33492742e+00]])
```

```
In [16]: type(X)
```

```
Out[16]: numpy.ndarray
```

```
In [17]: # Response/Target variable needs to a vector
# y must be a vector
y = df[['y']].to_numpy()
```

```
In [18]: # it is right now a array of dimensions 10x1
y.shape
```

```
Out[18]: (10, 1)
```

```
In [19]: y
```

```
Out[19]: array([[0],
 [0],
 [1],
 [1],
 [1],
 [1],
 [0],
 [0],
 [0],
 [1]])
```

```
In [20]: # Flatten to a single dimension array of 10 elements
y = y.ravel()
```

```
In [26]: y
```

```
Out[26]: array([0, 0, 1, 1, 1, 1, 0, 0, 0, 1])
```

```
In [23]: def write_recordio_file (filename, x, y=None):
          with open(filename, 'wb') as f:
              smac.write_numpy_to_dense_tensor(f, x, y)
```

```
In [24]: def read_recordio_file (filename, recordsToPrint = 10):
          with open(filename, 'rb') as f:
              record = smac.read_records(f)
              for i, r in enumerate(record):
                  if i >= recordsToPrint:
                      break
                  print ("record: {}".format(i))
                  print(r)
```

```
In [25]: write_recordio_file('demo_file.recordio',X,y)
```

```
In [27]: df.head(3)
```

```
Out[27]:
```

	x1	x2	x3	y
0	0.221993	153	2.041547	0
1	0.870732	180	1.190954	0
2	0.206719	127	8.779031	1

```
In [28]: read_recordio_file('demo_file.recordio',3)
```

```
record: 0
features {
  key: "values"
  value {
    float64_tensor {
      values: 0.22199317108973948
      values: 153.0
      values: 2.0415474783059215
    }
  }
}
label {
  key: "values"
  value {
    int32_tensor {
      values: 0
    }
  }
}
```

```
record: 1
features {
  key: "values"
  value {
    float64_tensor {
      values: 0.8707323061773764
      values: 180.0
      values: 1.1909535747826039
    }
  }
}
label {
  key: "values"
  value {
    int32_tensor {
      values: 0
    }
  }
}
```

```
record: 2
features {
  key: "values"
  value {
    float64_tensor {
      values: 0.20671915533942642
      values: 127.0
      values: 8.779030712603621
    }
  }
}
label {
  key: "values"
  value {
    int32_tensor {
      values: 1
    }
  }
}
```

```
    }  
  }  
}
```

```
In [29]: write_to_s3('demo_file.recordio', s3_bucket_name, 'data_format/demo_file.recordio')
```

```
In [30]: download_from_s3('demo_file_from_s3.recordio', s3_bucket_name, 'data_format/demo_file
```

```
In [ ]:
```