

# Train a model with Iris data using XGBoost algorithm

Model is trained with XGBoost installed in notebook instance

In the later examples, we will train using SageMaker's XGBoost algorithm

```
In [1]: # Install xgboost in notebook instance.  
##### Command to install xgboost  
!pip install xgboost
```

Looking in indexes: <https://pypi.org/simple>, <https://pip.repos.neuron.amazonaws.com>  
Requirement already satisfied: xgboost in /home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages (1.7.6)  
Requirement already satisfied: numpy in /home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages (from xgboost) (1.22.3)  
Requirement already satisfied: scipy in /home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages (from xgboost) (1.10.1)

```
In [2]: import sys  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import itertools  
import xgboost as xgb  
  
from sklearn import preprocessing  
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [3]: column_list_file = 'iris_train_column_list.txt'  
train_file = 'iris_train.csv'  
validation_file = 'iris_validation.csv'
```

```
In [4]: columns = ''  
with open(column_list_file, 'r') as f:  
    columns = f.read().split(',')
```

```
In [5]: columns
```

Out[5]: ['encoded\_class', 'sepal\_length', 'sepal\_width', 'petal\_length', 'petal\_width']

```
In [6]: # Encode Class Labels to integers
# Labeled Classes
labels=[0,1,2]
classes = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
le = preprocessing.LabelEncoder()
le.fit(classes)
```

Out[6]: ▾ LabelEncoder  
LabelEncoder()

```
In [7]: # Specify the column names as the file does not have column header
df_train = pd.read_csv(train_file,names=columns)
df_validation = pd.read_csv(validation_file,names=columns)
```

In [8]: df\_train.head()

Out[8]:

	encoded_class	sepal_length	sepal_width	petal_length	petal_width
0	1	5.8	2.7	3.9	1.2
1	2	6.1	2.6	5.6	1.4
2	2	5.8	2.8	5.1	2.4
3	0	4.4	3.2	1.3	0.2
4	2	7.2	3.6	6.1	2.5

In [9]: df\_validation.head()

Out[9]:

	encoded_class	sepal_length	sepal_width	petal_length	petal_width
0	1	5.8	2.7	4.1	1.0
1	0	4.8	3.4	1.6	0.2
2	1	6.0	2.2	4.0	1.0
3	2	6.4	3.1	5.5	1.8
4	2	6.7	2.5	5.8	1.8

```
In [10]: X_train = df_train.iloc[:,1:] # Features: 1st column onwards
y_train = df_train.iloc[:,0].ravel() # Target: 0th column

X_validation = df_validation.iloc[:,1:]
y_validation = df_validation.iloc[:,0].ravel()
```

```
In [11]: # Launch a classifier
# XGBoost Training Parameter Reference:
# https://xgboost.readthedocs.io/en/latest/parameter.html

classifier = xgb.XGBClassifier(objective="multi:softmax",
                               num_class=3,
                               n_estimators=100)
```

```
In [12]: classifier
```

Out[12]:

XGBClassifier

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytrees=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
```

```
In [13]: classifier.fit(X_train,
                        y_train,
                        eval_set = [(X_train, y_train), (X_validation, y_validation)],
                        eval_metric=['mlogloss'])
```

[0]	validation_0-mlogloss:0.73876	validation_1-mlogloss:0.74994
[1]	validation_0-mlogloss:0.52787	validation_1-mlogloss:0.55401
[2]	validation_0-mlogloss:0.38959	validation_1-mlogloss:0.42612
[3]	validation_0-mlogloss:0.29429	validation_1-mlogloss:0.34328
[4]	validation_0-mlogloss:0.22736	validation_1-mlogloss:0.29000
[5]	validation_0-mlogloss:0.17920	validation_1-mlogloss:0.24961
[6]	validation_0-mlogloss:0.14403	validation_1-mlogloss:0.22234
[7]	validation_0-mlogloss:0.11664	validation_1-mlogloss:0.20338
[8]	validation_0-mlogloss:0.09668	validation_1-mlogloss:0.18999
[9]	validation_0-mlogloss:0.08128	validation_1-mlogloss:0.18190
[10]	validation_0-mlogloss:0.06783	validation_1-mlogloss:0.17996
[11]	validation_0-mlogloss:0.05794	validation_1-mlogloss:0.18029
[12]	validation_0-mlogloss:0.05011	validation_1-mlogloss:0.18306
[13]	validation_0-mlogloss:0.04428	validation_1-mlogloss:0.18471
[14]	validation_0-mlogloss:0.03993	validation_1-mlogloss:0.18693
[15]	validation_0-mlogloss:0.03615	validation_1-mlogloss:0.18553
[16]	validation_0-mlogloss:0.03310	validation_1-mlogloss:0.18571
[17]	validation_0-mlogloss:0.03065	validation_1-mlogloss:0.18615
[18]	validation_0-mlogloss:0.02874	validation_1-mlogloss:0.18930
[19]	validation_0-mlogloss:0.02739	validation_1-mlogloss:0.18989
[20]	validation_0-mlogloss:0.02639	validation_1-mlogloss:0.19251
[21]	validation_0-mlogloss:0.02583	validation_1-mlogloss:0.19567
[22]	validation_0-mlogloss:0.02513	validation_1-mlogloss:0.19760
[23]	validation_0-mlogloss:0.02444	validation_1-mlogloss:0.19690
[24]	validation_0-mlogloss:0.02398	validation_1-mlogloss:0.19946
[25]	validation_0-mlogloss:0.02340	validation_1-mlogloss:0.20132
[26]	validation_0-mlogloss:0.02287	validation_1-mlogloss:0.20281
[27]	validation_0-mlogloss:0.02250	validation_1-mlogloss:0.20464
[28]	validation_0-mlogloss:0.02217	validation_1-mlogloss:0.20638
[29]	validation_0-mlogloss:0.02185	validation_1-mlogloss:0.20661
[30]	validation_0-mlogloss:0.02150	validation_1-mlogloss:0.20768
[31]	validation_0-mlogloss:0.02122	validation_1-mlogloss:0.20791
[32]	validation_0-mlogloss:0.02091	validation_1-mlogloss:0.21019
[33]	validation_0-mlogloss:0.02064	validation_1-mlogloss:0.21058
[34]	validation_0-mlogloss:0.02038	validation_1-mlogloss:0.21031
[35]	validation_0-mlogloss:0.02010	validation_1-mlogloss:0.21248
[36]	validation_0-mlogloss:0.01989	validation_1-mlogloss:0.21323
[37]	validation_0-mlogloss:0.01964	validation_1-mlogloss:0.21301
[38]	validation_0-mlogloss:0.01945	validation_1-mlogloss:0.21486
[39]	validation_0-mlogloss:0.01927	validation_1-mlogloss:0.21497
[40]	validation_0-mlogloss:0.01907	validation_1-mlogloss:0.21486
[41]	validation_0-mlogloss:0.01890	validation_1-mlogloss:0.21675

[42]	validation_0-mlogloss:0.01872	validation_1-mlogloss:0.21662
[43]	validation_0-mlogloss:0.01856	validation_1-mlogloss:0.21786
[44]	validation_0-mlogloss:0.01839	validation_1-mlogloss:0.21843
[45]	validation_0-mlogloss:0.01824	validation_1-mlogloss:0.21844
[46]	validation_0-mlogloss:0.01809	validation_1-mlogloss:0.21965
[47]	validation_0-mlogloss:0.01794	validation_1-mlogloss:0.21966
[48]	validation_0-mlogloss:0.01780	validation_1-mlogloss:0.22028
[49]	validation_0-mlogloss:0.01766	validation_1-mlogloss:0.22134
[50]	validation_0-mlogloss:0.01752	validation_1-mlogloss:0.22137
[51]	validation_0-mlogloss:0.01740	validation_1-mlogloss:0.22236
[52]	validation_0-mlogloss:0.01727	validation_1-mlogloss:0.22295
[53]	validation_0-mlogloss:0.01715	validation_1-mlogloss:0.22300
[54]	validation_0-mlogloss:0.01703	validation_1-mlogloss:0.22396
[55]	validation_0-mlogloss:0.01692	validation_1-mlogloss:0.22390
[56]	validation_0-mlogloss:0.01681	validation_1-mlogloss:0.22542
[57]	validation_0-mlogloss:0.01670	validation_1-mlogloss:0.22549
[58]	validation_0-mlogloss:0.01659	validation_1-mlogloss:0.22555
[59]	validation_0-mlogloss:0.01650	validation_1-mlogloss:0.22563
[60]	validation_0-mlogloss:0.01641	validation_1-mlogloss:0.22627
[61]	validation_0-mlogloss:0.01632	validation_1-mlogloss:0.22637
[62]	validation_0-mlogloss:0.01623	validation_1-mlogloss:0.22647
[63]	validation_0-mlogloss:0.01620	validation_1-mlogloss:0.22713
[64]	validation_0-mlogloss:0.01617	validation_1-mlogloss:0.22724
[65]	validation_0-mlogloss:0.01614	validation_1-mlogloss:0.22734
[66]	validation_0-mlogloss:0.01611	validation_1-mlogloss:0.22797
[67]	validation_0-mlogloss:0.01609	validation_1-mlogloss:0.22796
[68]	validation_0-mlogloss:0.01606	validation_1-mlogloss:0.22856
[69]	validation_0-mlogloss:0.01604	validation_1-mlogloss:0.22867
[70]	validation_0-mlogloss:0.01602	validation_1-mlogloss:0.22867
[71]	validation_0-mlogloss:0.01599	validation_1-mlogloss:0.22926
[72]	validation_0-mlogloss:0.01597	validation_1-mlogloss:0.22926
[73]	validation_0-mlogloss:0.01595	validation_1-mlogloss:0.22983

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `eval\_metric` in `fit` method is deprecated for better compatibility with scikit-learn, use `eval\_metric` in constructor or `set\_params` instead.

warnings.warn(

```

[74] validation_0-mlogloss:0.01592 validation_1-mlogloss:0.22994
[75] validation_0-mlogloss:0.01590 validation_1-mlogloss:0.22994
[76] validation_0-mlogloss:0.01588 validation_1-mlogloss:0.23050
[77] validation_0-mlogloss:0.01586 validation_1-mlogloss:0.23051
[78] validation_0-mlogloss:0.01584 validation_1-mlogloss:0.23106
[79] validation_0-mlogloss:0.01582 validation_1-mlogloss:0.23116
[80] validation_0-mlogloss:0.01580 validation_1-mlogloss:0.23117
[81] validation_0-mlogloss:0.01578 validation_1-mlogloss:0.23171
[82] validation_0-mlogloss:0.01576 validation_1-mlogloss:0.23172
[83] validation_0-mlogloss:0.01574 validation_1-mlogloss:0.23224
[84] validation_0-mlogloss:0.01573 validation_1-mlogloss:0.23225
[85] validation_0-mlogloss:0.01571 validation_1-mlogloss:0.23277
[86] validation_0-mlogloss:0.01569 validation_1-mlogloss:0.23287
[87] validation_0-mlogloss:0.01567 validation_1-mlogloss:0.23288
[88] validation_0-mlogloss:0.01566 validation_1-mlogloss:0.23339
[89] validation_0-mlogloss:0.01564 validation_1-mlogloss:0.23340
[90] validation_0-mlogloss:0.01562 validation_1-mlogloss:0.23390
[91] validation_0-mlogloss:0.01561 validation_1-mlogloss:0.23390
[92] validation_0-mlogloss:0.01559 validation_1-mlogloss:0.23439
[93] validation_0-mlogloss:0.01558 validation_1-mlogloss:0.23449
[94] validation_0-mlogloss:0.01556 validation_1-mlogloss:0.23450
[95] validation_0-mlogloss:0.01555 validation_1-mlogloss:0.23498
[96] validation_0-mlogloss:0.01553 validation_1-mlogloss:0.23499
[97] validation_0-mlogloss:0.01552 validation_1-mlogloss:0.23546
[98] validation_0-mlogloss:0.01550 validation_1-mlogloss:0.23556
[99] validation_0-mlogloss:0.01549 validation_1-mlogloss:0.23556

```

Out[13]:

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,

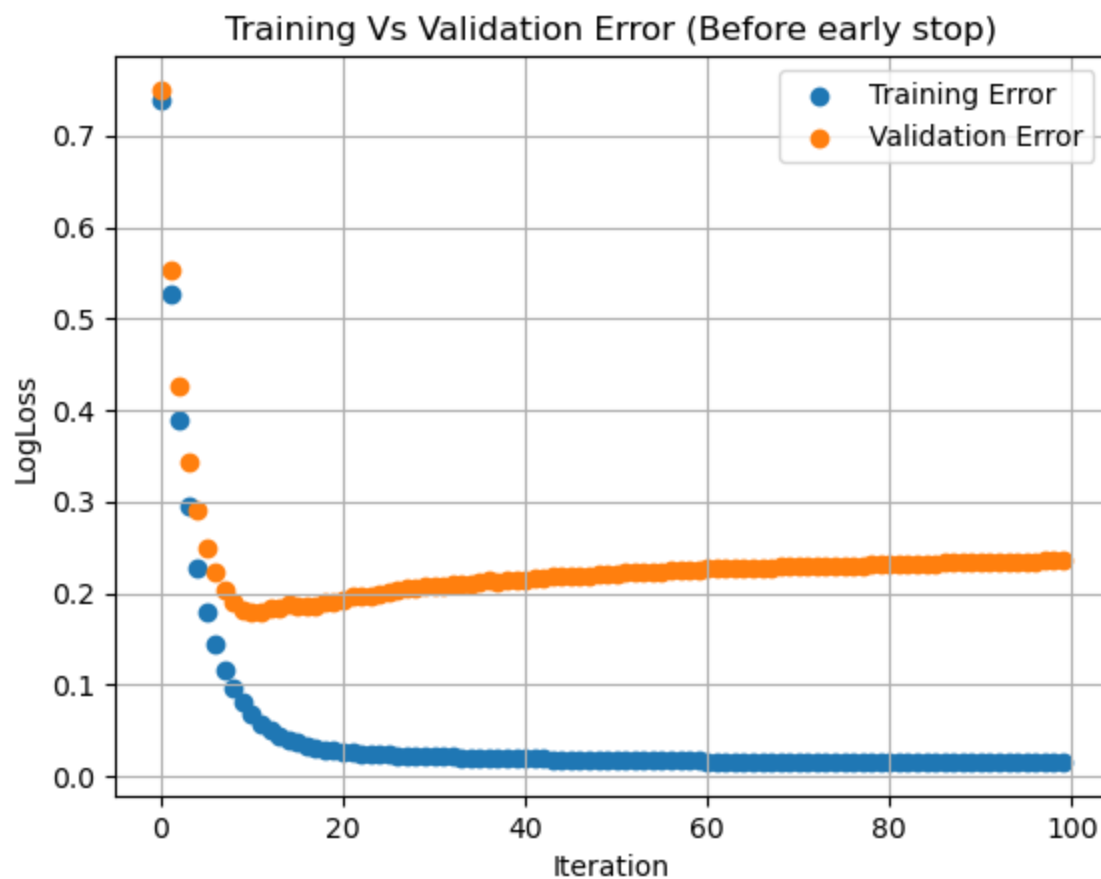
```

```
In [14]: eval_result_before_early_stop = classifier.eval_result()  
training_rounds = range(len(eval_result_before_early_stop['validation_0']['mlogloss']))  
print(f"training_rounds: {training_rounds}")
```

```
training_rounds: range(0, 100)
```

```
In [15]: plt.scatter(x=training_rounds,  
                    y=eval_result_before_early_stop[  
                        'validation_0']['mlogloss'],  
                    label='Training Error')  
plt.scatter(x=training_rounds,  
            y=eval_result_before_early_stop[  
                'validation_1']['mlogloss'], label='Validation Error')  
plt.grid(True)  
plt.xlabel('Iteration')  
plt.ylabel('LogLoss')  
plt.title('Training Vs Validation Error (Before early stop)')  
plt.legend()  
plt.show()
```





```
In [16]: classifier.fit(X_train,
                        y_train,
                        eval_set = [(X_train, y_train), (X_validation, y_validation)],
                        eval_metric=['mlogloss'],
                        early_stopping_rounds=10)

# early_stopping_rounds - needs to be passed in as a hyperparameter in SageMaker XGBoost implementation
# "The model trains until the validation score stops improving.
# Validation error needs to decrease at least every early_stopping_rounds to continue training.
# Amazon SageMaker hosting uses the best model for inference."
```

[0]	validation_0-mlogloss:0.73876	validation_1-mlogloss:0.74994
[1]	validation_0-mlogloss:0.52787	validation_1-mlogloss:0.55401
[2]	validation_0-mlogloss:0.38959	validation_1-mlogloss:0.42612
[3]	validation_0-mlogloss:0.29429	validation_1-mlogloss:0.34328
[4]	validation_0-mlogloss:0.22736	validation_1-mlogloss:0.29000
[5]	validation_0-mlogloss:0.17920	validation_1-mlogloss:0.24961
[6]	validation_0-mlogloss:0.14403	validation_1-mlogloss:0.22234
[7]	validation_0-mlogloss:0.11664	validation_1-mlogloss:0.20338
[8]	validation_0-mlogloss:0.09668	validation_1-mlogloss:0.18999
[9]	validation_0-mlogloss:0.08128	validation_1-mlogloss:0.18190
[10]	validation_0-mlogloss:0.06783	validation_1-mlogloss:0.17996
[11]	validation_0-mlogloss:0.05794	validation_1-mlogloss:0.18029
[12]	validation_0-mlogloss:0.05011	validation_1-mlogloss:0.18306
[13]	validation_0-mlogloss:0.04428	validation_1-mlogloss:0.18471
[14]	validation_0-mlogloss:0.03993	validation_1-mlogloss:0.18693
[15]	validation_0-mlogloss:0.03615	validation_1-mlogloss:0.18553
[16]	validation_0-mlogloss:0.03310	validation_1-mlogloss:0.18571
[17]	validation_0-mlogloss:0.03065	validation_1-mlogloss:0.18615
[18]	validation_0-mlogloss:0.02874	validation_1-mlogloss:0.18930
[19]	validation_0-mlogloss:0.02739	validation_1-mlogloss:0.18989

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `eval\_metric` in `fit` method is deprecated for better compatibility with scikit-learn, use `eval\_metric` in constructor or `set\_params` instead.

warnings.warn(

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `early\_stopping\_rounds` in `fit` method is deprecated for better compatibility with scikit-learn, use `early\_stopping\_rounds` in constructor or `set\_params` instead.

warnings.warn(

Out[16]:

```
XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
```


```
In [17]: eval_result = classifier.evals_result()
```

```
In [18]: training_rounds = range(len(eval_result['validation_0']['mlogloss']))
```


```
In [19]: print(training_rounds)

range(0, 21)
```

```
In [20]: plt.scatter(x=training_rounds,y=eval_result['validation_0']['mlogloss'],label='Training Error')
plt.scatter(x=training_rounds,y=eval_result['validation_1']['mlogloss'],label='Validation Error')
plt.grid(True)
plt.xlabel('Iteration')
plt.ylabel('LogLoss')
plt.title('Training Vs Validation Error')
plt.legend()
plt.show()
```



```
In [21]: xgb.plot_importance(classifier)
plt.show()
```



```
In [22]: df = pd.read_csv(validation_file,names=columns)
```

```
In [23]: df.head()
```

Out[23]:

	encoded_class	sepal_length	sepal_width	petal_length	petal_width
0	1	5.8	2.7	4.1	1.0
1	0	4.8	3.4	1.6	0.2
2	1	6.0	2.2	4.0	1.0
3	2	6.4	3.1	5.5	1.8
4	2	6.7	2.5	5.8	1.8

In [24]: `X_test = df.iloc[:,1:]`  
`print(X_test[:5])`

	sepal_length	sepal_width	petal_length	petal_width
0	5.8	2.7	4.1	1.0
1	4.8	3.4	1.6	0.2
2	6.0	2.2	4.0	1.0
3	6.4	3.1	5.5	1.8
4	6.7	2.5	5.8	1.8

In [25]: `result = classifier.predict(X_test)`

In [26]: `result[:5]`

Out[26]: `array([1, 0, 1, 2, 2], dtype=int32)`


In [27]: `df['predicted_class'] = result #le.inverse_transform(result)`  
`#DWB# = le.inverse_transform(result) #DWB# to get class names`

In [28]: `df.head()`

Out[28]:

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	predicted_class
0	1	5.8	2.7	4.1	1.0	1
1	0	4.8	3.4	1.6	0.2	0
2	1	6.0	2.2	4.0	1.0	1
3	2	6.4	3.1	5.5	1.8	2
4	2	6.7	2.5	5.8	1.8	2

```
In [29]: # Compare performance of Actual and Model 1 Prediction
plt.figure()
plt.scatter(df.index, df['encoded_class'], label='Actual')
plt.scatter(df.index, df['predicted_class'], label='Predicted', marker='^')
plt.legend(loc=4)
plt.yticks([0, 1, 2])
plt.xlabel('Sample')
plt.ylabel('Class')
plt.show()
```



```
In [30]: #DWB# After I make a copy of this DataFrame that Chandra had
#DWB#+ in the notebook
df_orig = df.copy(deep=True)

#DWB# Doing what was commented
df['pred_cls_decoded'] = le.inverse_transform(result)
df['decoded_cls'] = le.inverse_transform(df['encoded_class'])
```

```
In [31]: #DWB# Seeing result of doing what was commented
df.head()
```



Out[31]:

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	predicted_class	pred_cls_decoded	decoded_cls
0	1	5.8	2.7	4.1	1.0	1	Iris-versicolor	Iris-versicolor
1	0	4.8	3.4	1.6	0.2	0	Iris-setosa	Iris-setosa
2	1	6.0	2.2	4.0	1.0	1	Iris-versicolor	Iris-versicolor
3	2	6.4	3.1	5.5	1.8	2	Iris-virginica	Iris-virginica
4	2	6.7	2.5	5.8	1.8	2	Iris-virginica	Iris-virginica

In [32]: `print(df.head())`

```

    encoded_class  sepal_length  sepal_width  petal_length  petal_width
0              1           5.8           2.7           4.1           1.0 \
1              0           4.8           3.4           1.6           0.2
2              1           6.0           2.2           4.0           1.0
3              2           6.4           3.1           5.5           1.8
4              2           6.7           2.5           5.8           1.8

    predicted_class  pred_cls_decoded  decoded_cls
0              1  Iris-versicolor  Iris-versicolor
1              0   Iris-setosa    Iris-setosa
2              1  Iris-versicolor  Iris-versicolor
3              2  Iris-virginica  Iris-virginica
4              2  Iris-virginica  Iris-virginica

```

In [33]: `#DWB# Getting the data back to normal for the confusion matrix`  
`df['pred_cls_enc'] = result`  
`df['enc_cls_chk'] = le.transform(df['decoded_cls'])`

In [34]: `#DWB# Checking it worked.`  
`df.head()`

```
Out[34]:
```

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	predicted_class	pred_cls_decoded	decoded_cls	pred_cls_enc	enc_cls_chk
0	1	5.8	2.7	4.1	1.0	1	Iris-versicolor	Iris-versicolor	1	
1	0	4.8	3.4	1.6	0.2	0	Iris-setosa	Iris-setosa	0	
2	1	6.0	2.2	4.0	1.0	1	Iris-versicolor	Iris-versicolor	1	
3	2	6.4	3.1	5.5	1.8	2	Iris-virginica	Iris-virginica	2	
4	2	6.7	2.5	5.8	1.8	2	Iris-virginica	Iris-virginica	2	

```
In [35]: print(df.head())
```

```

  encoded_class  sepal_length  sepal_width  petal_length  petal_width
0             1           5.8           2.7           4.1           1.0 \
1             0           4.8           3.4           1.6           0.2
2             1           6.0           2.2           4.0           1.0
3             2           6.4           3.1           5.5           1.8
4             2           6.7           2.5           5.8           1.8

  predicted_class  pred_cls_decoded  decoded_cls  pred_cls_enc
0                1  Iris-versicolor  Iris-versicolor          1 \
1                0   Iris-setosa    Iris-setosa          0
2                1  Iris-versicolor  Iris-versicolor          1
3                2  Iris-virginica  Iris-virginica          2
4                2  Iris-virginica  Iris-virginica          2

  enc_cls_chk
0            1
1            0
2            1
3            2
4            2
```

```
In [36]: #DWB# I will actually get rid of those extra columns
#DWB#+ before starting the next section.
df.drop(columns=['pred_cls_decoded', 'decoded_cls',
                'pred_cls_enc', 'enc_cls_chk'], inplace=True)
```

```
In [37]: # Checking result
df.head()
```

```
Out[37]:
```

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	predicted_class
0	1	5.8	2.7	4.1	1.0	1
1	0	4.8	3.4	1.6	0.2	0
2	1	6.0	2.2	4.0	1.0	1
3	2	6.4	3.1	5.5	1.8	2
4	2	6.7	2.5	5.8	1.8	2

```
In [38]: # Quick, non-thorough check that we're back
df_orig.head()
```

```
Out[38]:
```

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	predicted_class
0	1	5.8	2.7	4.1	1.0	1
1	0	4.8	3.4	1.6	0.2	0
2	1	6.0	2.2	4.0	1.0	1
3	2	6.4	3.1	5.5	1.8	2
4	2	6.7	2.5	5.8	1.8	2

## Confusion Matrix

Confusion Matrix is a table that summarizes performance of classification model.

```
In [39]: # Reference:
# https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
```

```

"""
This function prints and plots the confusion matrix.
Normalization can be applied by setting `normalize=True`.
"""
if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
    #print("Normalized confusion matrix")
#else:
#    print('Confusion matrix, without normalization')

#print(cm)

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()

```

```

In [40]: # Compute confusion matrix
cnf_matrix = confusion_matrix(df['encoded_class'],
                             df['predicted_class'], labels=labels)

```

```

In [41]: cnf_matrix


```

```


Out[41]: array([[16,  0,  0],
                [ 0, 10,  1],
                [ 0,  1, 17]])

```

```
In [42]: # Plot confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=classes,
                      title='Confusion matrix - Count')
```



```
In [43]: # Plot confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=classes,
                      title='Confusion matrix - Count', normalize=True)
```



```
In [ ]: #DWB# Note that, with this being a multi-class classification problem,  
#DWB#+ weighted average f1-score is a good metric.  
print(classification_report(df['encoded_class'],  
                             df['predicted_class'],  
                             labels=labels,  
                             target_names=classes))
```

```
In [ ]:
```