

```
In [1]: import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
import os

import boto3
import re # python regex module
from sagemaker import get_execution_role
import sagemaker

# SDK 2 serializers and deserializers
from sagemaker.serializers import CSVSerializer
from sagemaker.deserializers import JSONDeserializer
```

XGBoost Cloud Prediction - Iris Classification

Invoke SageMaker Prediction Service

```
In [2]: # Acquire a realtime endpoint
endpoint_name = 'xgboost-iris-v1' #DWB# Checked from console - matches
predictor = sagemaker.predictor.Predictor (endpoint_name=endpoint_name)
```

```
In [3]: predictor.serializer = CSVSerializer()
```

```
In [4]: # Test predictive quality against data in validation file
df_all = pd.read_csv('iris_validation.csv',
                    names=['encoded_class', 'sepal_length', 'sepal_width', 'petal_length', 'petal_width'])
```

```
In [5]: df_all.head()
```

```
Out[5]:
```

	encoded_class	sepal_length	sepal_width	petal_length	petal_width
0	1	5.8	2.7	4.1	1.0
1	0	4.8	3.4	1.6	0.2
2	1	6.0	2.2	4.0	1.0
3	2	6.4	3.1	5.5	1.8
4	2	6.7	2.5	5.8	1.8

```
In [6]: df_all.columns
```

```
Out[6]: Index(['encoded_class', 'sepal_length', 'sepal_width', 'petal_length',
              'petal_width'],
              dtype='object')
```

```
In [7]: # Need to pass an array to the prediction
        # can pass a numpy array or a list of values [[19,1],[20,1]]
        # arr_test = df_all.as_matrix(['sepal_length', 'sepal_width', 'petal_length', 'petal_width'])
        arr_test = df_all[['sepal_length', 'sepal_width', 'petal_length', 'petal_width']].values
```

```
In [8]: type(arr_test)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: arr_test.shape
```

```
Out[9]: (45, 4)
```

```
In [10]: arr_test[:5]
```

```
Out[10]: array([[5.8, 2.7, 4.1, 1. ],
                [4.8, 3.4, 1.6, 0.2],
                [6. , 2.2, 4. , 1. ],
                [6.4, 3.1, 5.5, 1.8],
                [6.7, 2.5, 5.8, 1.8]])
```

```
In [11]: result = predictor.predict(arr_test[:2])
```

```
In [12]: arr_test.shape
```

```
Out[12]: (45, 4)
```

```
In [15]: #DWB# I think that repeat might not be on purpose;
#DWB#+ Let's check instead result.shape
try:
    print(str(result.shape))
except Exception as e:
    print('', file=sys.stderr)
    print("That didn't work.", file=sys.stderr)
    print(f"str(e) is: `{str(e)}'", file=sys.stderr)
    print('', file=sys.stderr)
finally:
    print("That tells us what we need to know.")
##endof: try/except/finally
```

That tells us what we need to know.

That didn't work.

str(e) is: `'bytes' object has no attribute 'shape'`

```
In [16]: result
```

```
Out[16]: b'1.0\n0.0\n'
```

```
In [22]: # For large number of predictions, we can split the input data and
# Query the prediction service.
# array_split is convenient to specify how many splits are needed

# Splitting using regular expression as xgboost 1-2-2 is returning
# predicted values with inconsistent delimiters (comma, newline or both)

# pattern looks for one or more of non-numeric characters
pattern = r'^0-9.+'

predictions = []
#DWB# added the next 2 lines
total_row_count = 0
n_columns_and_count = {}
```

```
for arr in np.array_split(arr_test,10):
    result = predictor.predict(arr)
    result = re.split(pattern,result.decode("utf-8"))
    print (arr.shape)
    #DWB# Here is what we can match up
    #DWB# <shape-consistency-check>
    this_chunk_shape = arr.shape
    this_row_count = this_chunk_shape[0]
    total_row_count += this_row_count
    this_col_count = this_chunk_shape[1]
    if this_col_count in n_columns_and_count:
        n_columns_and_count[this_col_count] += 1
    else:
        n_columns_and_count[this_col_count] = 1
    ##endof: if/else this_col_count in n_columns_and_count
    #DWB# </shape-consistency-check>
    predictions += [int(float(r)) for r in result if r != ""]

#DWB# It's me from here on out.

print()
print("# Looking at the chunks all together #")
print(f"The total number of rows is: {total_row_count}")
print("For each row, I counted the number of columns;")
print("here is the distribution of column counts.")
print(n_columns_and_count)
print()
print("Having inspected that, I can see that")
print("the shape of all the chunks combined is")
print("(45, 4), which matches our original")
print("arr_test.")
```

```
(5, 4)
(5, 4)
(5, 4)
(5, 4)
(5, 4)
(4, 4)
(4, 4)
(4, 4)
(4, 4)
(4, 4)
```

```
# Looking at the chunks all together #
The total number of rows is: 45
For each row, I counted the number of columns;
here is the distribution of column counts.
{4: 10}
```

```
Having inspected that, I can see that
the shape of all the chunks combined is
(45, 4), which matches our original
arr_test.
```

```
In [23]: len(predictions)
```

```
Out[23]: 45
```

```
In [24]: predictions[:5]
```

```
Out[24]: [1, 0, 1, 2, 2]
```

```
In [25]: from sklearn import preprocessing
le = preprocessing.LabelEncoder()
le.fit(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'])
```

```
Out[25]: ▼ LabelEncoder
LabelEncoder()
```

```
In [26]: df_all['class'] = le.inverse_transform(df_all.encoded_class)
```

```
In [27]: df_all['predicted_class']=le.inverse_transform(predictions)
```

```
In [28]: df_all.head()
```

```
Out[28]:
```

	encoded_class	sepal_length	sepal_width	petal_length	petal_width	class	predicted_class
0	1	5.8	2.7	4.1	1.0	Iris-versicolor	Iris-versicolor
1	0	4.8	3.4	1.6	0.2	Iris-setosa	Iris-setosa
2	1	6.0	2.2	4.0	1.0	Iris-versicolor	Iris-versicolor
3	2	6.4	3.1	5.5	1.8	Iris-virginica	Iris-virginica
4	2	6.7	2.5	5.8	1.8	Iris-virginica	Iris-virginica

```
In [29]: print('Confusion matrix - Actual versus Predicted')
pd.crosstab(df_all['class'], df_all['predicted_class'])
```

Confusion matrix - Actual versus Predicted

```
Out[29]: predicted_class Iris-setosa Iris-versicolor Iris-virginica
```

class			
Iris-setosa	16	0	0
Iris-versicolor	0	10	1
Iris-virginica	0	1	17

```
In [30]: import sklearn.metrics as metrics
print(metrics.classification_report(df_all['class'], df_all['predicted_class']))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	16
Iris-versicolor	0.91	0.91	0.91	11
Iris-virginica	0.94	0.94	0.94	18
accuracy			0.96	45
macro avg	0.95	0.95	0.95	45
weighted avg	0.96	0.96	0.96	45

```
In [31]: #DWB# Still in this second one there's no Endpoint-deletion Code.
#DWB#+ I will put some in, here.
#DWB#+ As Chandra wrote with the previous such code
# Delete Endpoint to prevent unnecessary charges
predictor.delete_endpoint()
```

```
In [32]: # I checked the list of endpoints from the AWS Console > Sagemaker ...
#+ and the endpoint that was there is gone.
```

```
In [ ]:
```