

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

## Diabetes Binary Classification Dataset

Input Features: 'preg\_count', 'glucose\_concentration', 'diastolic\_bp', 'triceps\_skin\_fold\_thickness', 'two\_hr\_serum\_insulin', 'bmi', 'diabetes\_pedi', 'age'

Target Feature: 'diabetes\_class'

Objective: Predict diabetes\_class for given input features

Data Source: <https://archive.ics.uci.edu/ml/datasets/pima+indians+diabetes>

```
In [2]: columns = ['diabetes_class', 'preg_count', 'glucose_concentration', 'diastolic_bp',
                  'triceps_skin_fold_thickness', 'two_hr_serum_insulin', 'bmi',
                  'diabetes_pedi', 'age']
```

```
In [3]: df = pd.read_csv('pima_indians_diabetes_all.csv')
```

```
In [4]: df.describe() # Note that most of the minima are zero. That's not good.
```

Out[4]:

	preg_count	glucose_concentration	diastolic_bp	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240000
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760000
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000

In [17]: `print(df.describe())`

	preg_count	glucose_concentration	diastolic_bp	
count	768.000000	768.000000	768.000000	\
mean	3.845052	120.894531	69.105469	
std	3.369578	31.972618	19.355807	
min	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	
50%	3.000000	117.000000	72.000000	
75%	6.000000	140.250000	80.000000	
max	17.000000	199.000000	122.000000	


  

	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	
count	768.000000	768.000000	768.000000	\
mean	20.536458	79.799479	31.992578	
std	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	27.300000	
50%	23.000000	30.500000	32.000000	
75%	32.000000	127.250000	36.600000	
max	99.000000	846.000000	67.100000	

	diabetes_pedi	age	diabetes_class
count	768.000000	768.000000	768.000000
mean	0.471876	33.240885	0.348958
std	0.331329	11.760232	0.476951
min	0.078000	21.000000	0.000000
25%	0.243750	24.000000	0.000000
50%	0.372500	29.000000	0.000000
75%	0.626250	41.000000	1.000000
max	2.420000	81.000000	1.000000

```
In [5]: df['glucose_concentration'].hist()
plt.show()
```



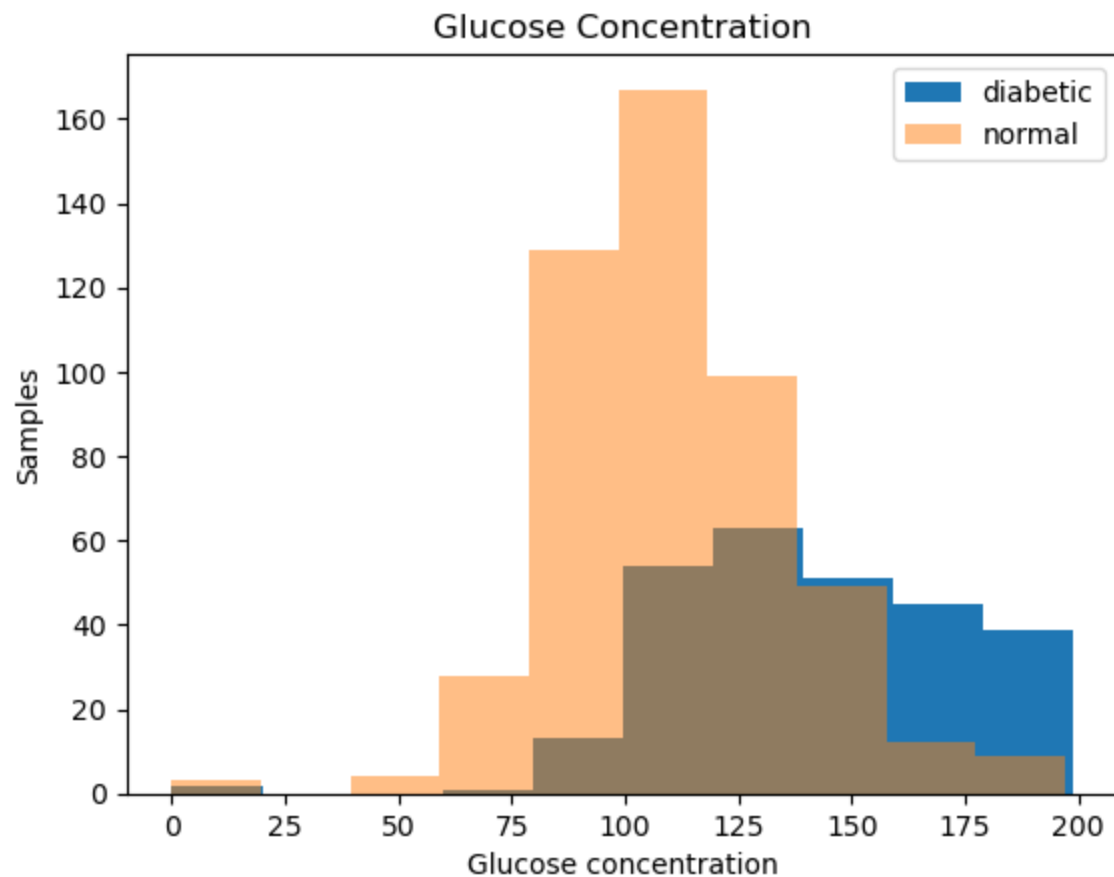
```
In [6]: df['diabetes_class'].value_counts()
```

```
Out[6]: diabetes_class  
0      500  
1      268  
Name: count, dtype: int64
```


```
In [7]: # Separate diabetic and normal samples  
diabetic = df.diabetes_class == 1  
normal = df.diabetes_class == 0
```

```
In [8]: # Glucose concentration histogram  
plt.hist(df[diabetic].glucose_concentration, label='diabetic')  
plt.hist(df[normal].glucose_concentration, alpha=0.5, label='normal')  
plt.title('Glucose Concentration')  
plt.xlabel('Glucose concentration')
```


```
plt.ylabel('Samples')  
plt.legend()  
plt.show()
```



```
In [9]: # BMI histogram  
plt.hist(df[diabetic].bmi,label='diabetic')  
plt.hist(df[normal].bmi,alpha=0.5,label='normal')  
plt.title('BMI')  
plt.xlabel('BMI')  
plt.ylabel('Samples')  
plt.legend()  
plt.show()
```



```
In [10]: # Age
plt.hist(df[diabetic].age, label='diabetic')
plt.hist(df[normal].age, alpha=0.5, label='normal')
plt.title('Age')
plt.xlabel('Age')
plt.ylabel('Samples')
plt.legend()
plt.show()
```



## Training and Validation Set

Target Variable as first column followed by input features:

'diabetes\_class', 'preg\_count', 'glucose\_concentration', 'diastolic\_bp', 'triceps\_skin\_fold\_thickness', 'two\_hr\_serum\_insulin', 'bmi',  
'diabetes\_pedi', 'age'

Training, Validation files do not have a column header

```
In [11]: # Training = 70% of the data  
# Validation = 30% of the data
```

```
# Randomize the dataset
np.random.seed(5)
l = list(df.index)
np.random.shuffle(l)
df = df.iloc[l]
```

```
In [12]: rows = df.shape[0]
train = int(.7 * rows)
test = rows - train
```

```
In [13]: rows, train, test
```

```
Out[13]: (768, 537, 231)
```

```
In [14]: # Write Training Set
df[:train].to_csv('diabetes_train.csv'
                  ,index=False,index_label='Row',header=False
                  ,columns=columns)
```

```
In [15]: # Write Validation Set
df[train:].to_csv('diabetes_validation.csv'
                  ,index=False,index_label='Row',header=False
                  ,columns=columns)
```

```
In [16]: # Write Column List
with open('diabetes_train_column_list.txt','w') as f:
    f.write(','.join(columns))
```

```
In [ ]:
```