

Diabetes dataset

Predict if a person is at risk of developing diabetes

Binary Classification problem - XGBoost

```
In [1]: # Install xgboost in notebook instance.  
##### Command to install xgboost  
!pip install xgboost
```

Looking in indexes: <https://pypi.org/simple>, <https://pip.repos.neuron.amazonaws.com>

Collecting xgboost

Downloading xgboost-1.7.6-py3-none-manylinux2014_x86_64.whl (200.3 MB)

200.3/200.3 MB 3.1 MB/s eta 0:00:0000:0100:01

Requirement already satisfied: numpy in /home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages (from xgboost) (1.22.3)

Requirement already satisfied: scipy in /home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages (from xgboost) (1.10.1)

Installing collected packages: xgboost

Successfully installed xgboost-1.7.6

```
In [2]: import sys  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import itertools  
  
import xgboost as xgb  
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [3]: column_list_file = 'diabetes_train_column_list.txt'  
train_file = 'diabetes_train.csv'  
validation_file = 'diabetes_validation.csv'
```

```
In [4]: columns = ''  
with open(column_list_file, 'r') as f:  
    columns = f.read().split(',')
```

In [5]: columns

```
Out[5]: ['diabetes_class',
        'preg_count',
        'glucose_concentration',
        'diastolic_bp',
        'triceps_skin_fold_thickness',
        'two_hr_serum_insulin',
        'bmi',
        'diabetes_pedi',
        'age']
```

```
In [6]: # Specify the column names as the file does not have column header
df_train = pd.read_csv(train_file,names=columns)
df_validation = pd.read_csv(validation_file,names=columns)
```

In [7]: df_train.head()

```
Out[7]:
```

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	0	6	92	62	32	126	32.0	0.085	4
1	0	5	132	80	0	0	26.8	0.186	6
2	0	3	106	72	0	0	25.8	0.207	2
3	0	4	99	68	38	0	32.8	0.145	3
4	0	4	96	56	17	49	20.8	0.340	2

In [8]: print(df_train.head())

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	
0	0	6	92	62	\
1	0	5	132	80	
2	0	3	106	72	
3	0	4	99	68	
4	0	4	96	56	

	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	32	126	32.0	0.085	46
1	0	0	26.8	0.186	69
2	0	0	25.8	0.207	27
3	38	0	32.8	0.145	33
4	17	49	20.8	0.340	26

In [9]: `df_validation.head()`

Out[9]:

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	0	1	130	70	13	105	25.9	0.472	2
1	1	8	133	72	0	0	32.9	0.270	3
2	0	0	137	68	14	148	24.8	0.143	2
3	0	2	88	74	19	53	29.0	0.229	2
4	1	9	130	70	0	0	34.2	0.652	4

In [10]: `print(df_validation.head())`

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	
0	0	1	130	70	\
1	1	8	133	72	
2	0	0	137	68	
3	0	2	88	74	
4	1	9	130	70	

	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	13	105	25.9	0.472	22
1	0	0	32.9	0.270	39
2	14	148	24.8	0.143	21
3	19	53	29.0	0.229	22
4	0	0	34.2	0.652	45

```
In [11]: X_train = df_train.iloc[:,1:] # Features: 1st column onwards
y_train = df_train.iloc[:,0].ravel() # Target: 0th column

X_validation = df_validation.iloc[:,1:]
y_validation = df_validation.iloc[:,0].ravel()
```

```
In [12]: # Launch a classifier
# XGBoost Training Parameter Reference:
# https://xgboost.readthedocs.io/en/latest/parameter.html
classifier = xgb.XGBClassifier (objective="binary:logistic")
```

```
In [13]: classifier
```

```
Out[13]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
```

In [41]: `print(classifier)`

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...)
```

In [14]: `classifier.fit(X_train,
 y_train,
 eval_set = [(X_train, y_train), (X_validation, y_validation)],
 eval_metric=['logloss'],
 early_stopping_rounds=10)`

```
[0]    validation_0-logloss:0.55122    validation_1-logloss:0.60755
[1]    validation_0-logloss:0.46796    validation_1-logloss:0.57246
[2]    validation_0-logloss:0.40847    validation_1-logloss:0.54856
[3]    validation_0-logloss:0.36336    validation_1-logloss:0.53960
[4]    validation_0-logloss:0.32749    validation_1-logloss:0.53086
[5]    validation_0-logloss:0.30366    validation_1-logloss:0.52163
[6]    validation_0-logloss:0.28429    validation_1-logloss:0.52666
[7]    validation_0-logloss:0.26441    validation_1-logloss:0.52752
[8]    validation_0-logloss:0.24931    validation_1-logloss:0.52957
[9]    validation_0-logloss:0.23460    validation_1-logloss:0.52743
[10]   validation_0-logloss:0.22422    validation_1-logloss:0.53384
[11]   validation_0-logloss:0.21326    validation_1-logloss:0.53555
[12]   validation_0-logloss:0.20211    validation_1-logloss:0.54085
[13]   validation_0-logloss:0.19681    validation_1-logloss:0.54294
[14]   validation_0-logloss:0.19060    validation_1-logloss:0.54895
```

```

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `eval_metric`
in `fit` method is deprecated for better compatibility with scikit-learn, use `eval_metric` in constructor or `set_
params` instead.
  warnings.warn(
/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `early_stopp
ing_rounds` in `fit` method is deprecated for better compatibility with scikit-learn, use `early_stopping_rounds` in
constructor or `set_params` instead.
  warnings.warn(

```

Out[14]:

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,

```


```
In [15]: eval_result = classifier.evals_result()
```

```
In [16]: training_rounds = range(len(eval_result['validation_0']['logloss']))
```


```
In [17]: print(training_rounds)
```

```
range(0, 16)
```

```
In [18]: plt.scatter(x=training_rounds,y=eval_result['validation_0']['logloss'],label='Training Error')
plt.scatter(x=training_rounds,y=eval_result['validation_1']['logloss'],label='Validation Error')
plt.grid(True)
plt.xlabel('Iteration')
plt.ylabel('LogLoss')
plt.title('Training Vs Validation Error')
plt.legend()
plt.show()
```



```
In [19]: xgb.plot_importance(classifier)
plt.show()
```



```
In [20]: df = pd.read_csv(validation_file,names=columns)
```

```
In [21]: df.head()
```



```
Out[21]:
```

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	0	1	130	70	13	105	25.9	0.472	22
1	1	8	133	72	0	0	32.9	0.270	39
2	0	0	137	68	14	148	24.8	0.143	21
3	0	2	88	74	19	53	29.0	0.229	22
4	1	9	130	70	0	0	34.2	0.652	45

```
In [22]: print(df.head())
```

```

   diabetes_class  preg_count  glucose_concentration  diastolic_bp
0              0           1             130           70
1              1           8             133           72
2              0           0             137           68
3              0           2              88           74
4              1           9             130           70

   triceps_skin_fold_thickness  two_hr_serum_insulin  bmi  diabetes_pedi  age
0                        13             105  25.9           0.472  22
1                        0               0  32.9           0.270  39
2                        14             148  24.8           0.143  21
3                        19              53  29.0           0.229  22
4                        0               0  34.2           0.652  45

```

```
In [23]: X_test = df.iloc[:,1:]
```

```
In [24]: result = classifier.predict(X_test)
```

```
In [25]: result[:5]
```

```
Out[25]: array([0, 0, 0, 0, 1])
```

```
In [26]: df['predicted_class'] = result
```

```
In [27]: df.head()
```

Out[27]:

	diabetes_class	preg_count	glucose_concentration	diastolic_bp	triceps_skin_fold_thickness	two_hr_serum_insulin	bmi	diabetes_pedi	age
0	0	1	130	70	13	105	25.9	0.472	22
1	1	8	133	72	0	0	32.9	0.270	39
2	0	0	137	68	14	148	24.8	0.143	21
3	0	2	88	74	19	53	29.0	0.229	22
4	1	9	130	70	0	0	34.2	0.652	45

In [28]: `print(df.head())`

```

diabetes_class  preg_count  glucose_concentration  diastolic_bp
0              0          1                   130           70 \
1              1          8                   133           72
2              0          0                   137           68
3              0          2                   88            74
4              1          9                   130           70

triceps_skin_fold_thickness  two_hr_serum_insulin  bmi  diabetes_pedi
0                          13                   105  25.9         0.472 \
1                          0                    0   32.9         0.270
2                          14                   148  24.8         0.143
3                          19                    53  29.0         0.229
4                          0                    0   34.2         0.652

age  predicted_class
0   22              0
1   39              0
2   21              0
3   22              0
4   45              1

```

Binary Classifier Metrics

In [29]: `# Reference: https://scikit-learn.org/stable/modules/model_evaluation.html
Explicitly stating labels. Pass=1, Fail=0
def true_positive(y_true, y_pred):`

```

    return confusion_matrix(y_true, y_pred, labels=[1,0])[0, 0]

def true_negative(y_true, y_pred):
    return confusion_matrix(y_true,y_pred,labels=[1,0])[1, 1]

def false_positive(y_true, y_pred):
    return confusion_matrix(y_true, y_pred,labels=[1,0])[1, 0]

def false_negative(y_true, y_pred):
    return confusion_matrix(y_true, y_pred,labels=[1,0])[0, 1]

```

```

In [30]: # Compute Binary Classifier Metrics
# Returns a dictionary {"MetricName":Value,...}

def binary_classifier_metrics(y_true, y_pred):
    metrics = {}

    # References:
    # https://docs.aws.amazon.com/machine-learning/latest/dg/binary-classification.html
    # https://en.wikipedia.org/wiki/Confusion\_matrix

    # Definition:
    # true positive = tp = how many samples were correctly classified as positive (count)
    # true negative = tn = how many samples were correctly classified as negative (count)
    # false positive = fp = how many negative samples were mis-classified as positive (count)
    # false_negative = fn = how many positive samples were mis-classified as negative (count)

    # positive = number of positive samples (count)
    #           = true positive + false negative
    # negative = number of negative samples (count)
    #           = true negative + false positive

    tp = true_positive(y_true, y_pred)
    tn = true_negative(y_true, y_pred)
    fp = false_positive(y_true, y_pred)
    fn = false_negative(y_true, y_pred)

    positive = tp + fn
    negative = tn + fp

    metrics['TruePositive'] = tp
    metrics['TrueNegative'] = tn

```

```
metrics['FalsePositive'] = fp
metrics['FalseNegative'] = fn

metrics['Positive'] = positive
metrics['Negative'] = negative

# True Positive Rate (TPR, Recall) = true positive/positive
# How many positives were correctly classified? (fraction)
# Recall value closer to 1 is better. closer to 0 is worse
if tp == 0:
    recall = 0
else:
    recall = tp/positive

metrics['Recall'] = recall

# True Negative Rate = True Negative/negative
# How many negatives were correctly classified? (fraction)
# True Negative Rate value closer to 1 is better. closer to 0 is worse
if tn == 0:
    tnr = 0
else:
    tnr = tn/(negative)
metrics['TrueNegativeRate'] = tnr

# Precision = True Positive/(True Positive + False Positive)
# How many positives classified by the algorithm are really positives? (fraction)
# Precision value closer to 1 is better. closer to 0 is worse
if tp == 0:
    precision = 0
else:
    precision = tp/(tp + fp)
metrics['Precision'] = precision

# Accuracy = (True Positive + True Negative)/(total positive + total negative)
# How many positives and negatives were correctly classified? (fraction)
# Accuracy value closer to 1 is better. closer to 0 is worse
accuracy = (tp + tn)/(positive + negative)
metrics['Accuracy'] = accuracy

# False Positive Rate (FPR, False Alarm) = False Positive/(total negative)
# How many negatives were mis-classified as positives (fraction)
```

```

# False Positive Rate value closer to 0 is better. closer to 1 is worse
if fp == 0:
    fpr = 0
else:
    fpr = fp/(negative)
metrics['FalsePositiveRate'] = fpr

# False Negative Rate (FNR, Misses) = False Negative/(total Positive)
# How many positives were mis-classified as negative (fraction)
# False Negative Rate value closer to 0 is better. closer to 1 is worse
fnr = fn/(positive)
metrics['FalseNegativeRate'] = fnr

# F1 Score = harmonic mean of Precision and Recall
# F1 Score closer to 1 is better. Closer to 0 is worse.
if precision == 0 or recall == 0:
    f1 = 0
else:
    f1 = 2*precision*recall/(precision+recall)

metrics['F1'] = f1

return metrics

```

```

In [31]: # Reference:
# https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        #print("Normalized confusion matrix")
    #else:
        #    print('Confusion matrix, without normalization')

    #print(cm)

```

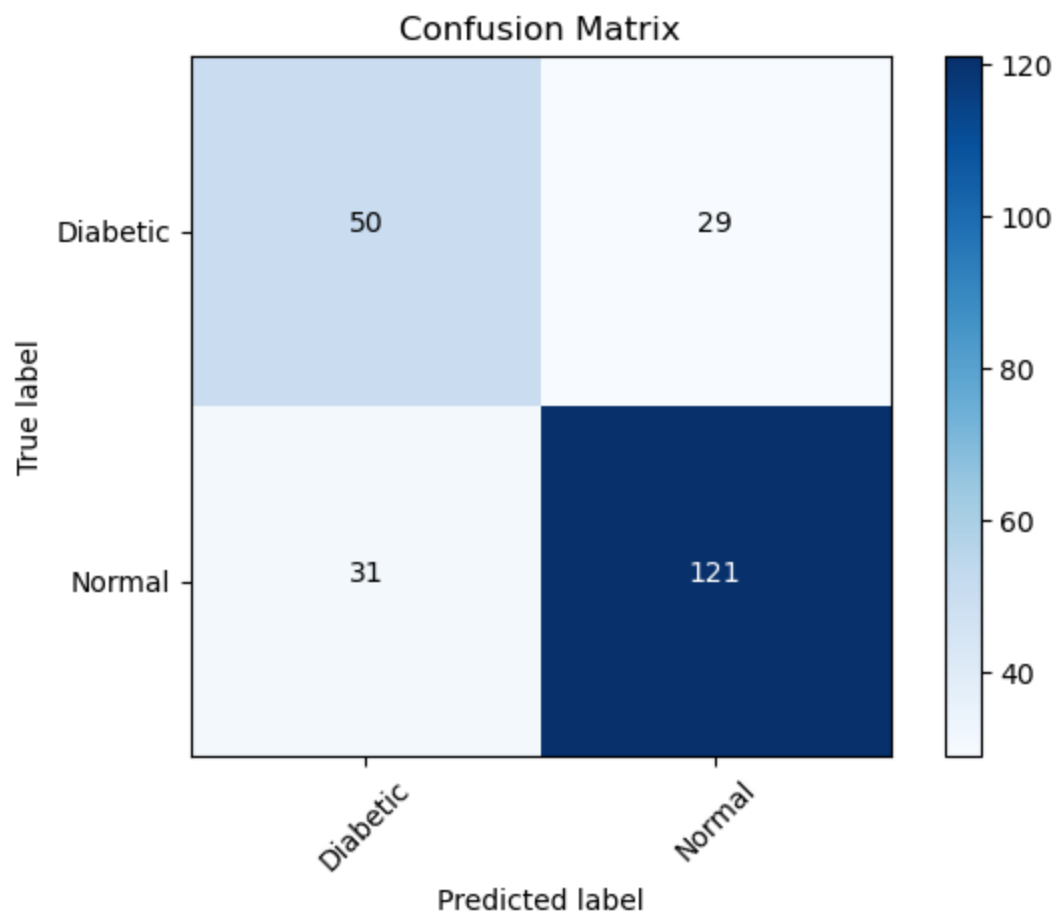
```
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

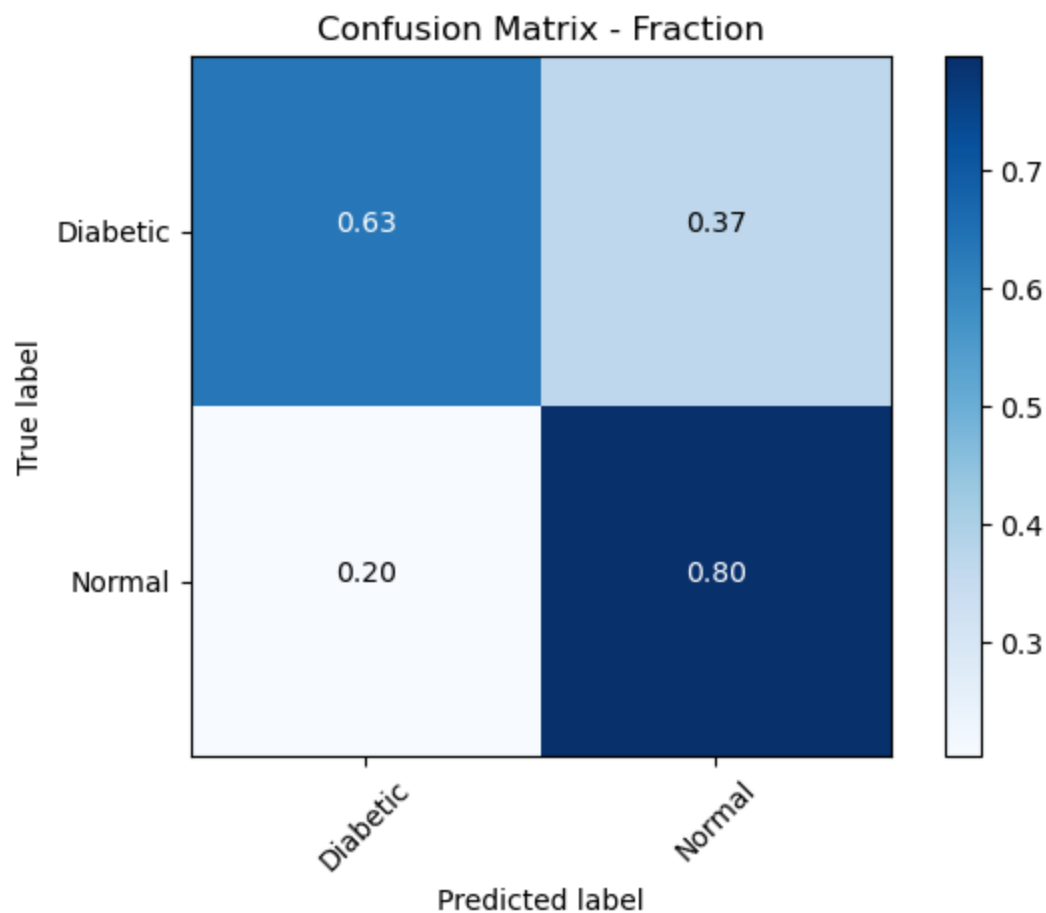
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
```

```
In [32]: # Compute confusion matrix
cnf_matrix = confusion_matrix(df['diabetes_class'], df['predicted_class'], labels=[1,0])
```

```
In [33]: # Plot confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['Diabetic', 'Normal'],
                      title='Confusion Matrix')
```



```
In [34]: # Plot confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['Diabetic', 'Normal'],
                      title='Confusion Matrix - Fraction', normalize=True)
```



```
In [35]: metrics = [binary_classifier_metrics(df['diabetes_class'], df['predicted_class'])]
df_metrics=pd.DataFrame.from_dict(metrics)
df_metrics.index = ['Model']
```

```
In [36]: df_metrics
```

```
Out[36]:
```

	TruePositive	TrueNegative	FalsePositive	FalseNegative	Positive	Negative	Recall	TrueNegativeRate	Precision	Accuracy	False
Model	50	121	31	29	79	152	0.632911	0.796053	0.617284	0.74026	

```
In [37]: print(df_metrics)
```


	TruePositive	TrueNegative	FalsePositive	FalseNegative	Positive
Model	50	121	31	29	79 \

	Negative	Recall	TrueNegativeRate	Precision	Accuracy
Model	152	0.632911	0.796053	0.617284	0.74026 \

	FalsePositiveRate	FalseNegativeRate	F1
Model	0.203947	0.367089	0.625

```
In [38]: print('Counts')
print(df_metrics[['TruePositive',
                  'FalseNegative',
                  'FalsePositive',
                  'TrueNegative']].round(2))

print()
print('Fractions')
print(df_metrics[['Recall',
                  'FalseNegativeRate',
                  'FalsePositiveRate',
                  'TrueNegativeRate']].round(2))

print()

print(df_metrics[['Precision',
                  'Accuracy',
                  'F1']].round(2))
```

Counts

	TruePositive	FalseNegative	FalsePositive	TrueNegative
Model	50	29	31	121

Fractions

	Recall	FalseNegativeRate	FalsePositiveRate	TrueNegativeRate
Model	0.63	0.37	0.2	0.8

	Precision	Accuracy	F1
Model	0.62	0.74	0.63

```
In [39]: print(classification_report(
    df['diabetes_class'],
    df['predicted_class'],
    labels=[1,0],
    target_names=['Diabetic', 'Normal']))
```

	precision	recall	f1-score	support
Diabetic	0.62	0.63	0.63	79
Normal	0.81	0.80	0.80	152
accuracy			0.74	231
macro avg	0.71	0.71	0.71	231
weighted avg	0.74	0.74	0.74	231

```
In [40]: # Yeah, not so good. Those dang zeros. We're going to fix it.
```

```
In [ ]:
```